

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error,
mean_squared_error, r2_score
```

```
In [2]: # Load dataset
df = pd.read_csv('BD_Weather.csv')
```

In [3]: # EDA

```
print("Dataset shape:", df.shape)
print("\nMissing values:\n", df.isnull().sum())
print("\nSummary statistics:\n", df.describe())
```

Dataset shape: (76685, 8)

Missing values:

Station	0
Year	0
Month	0
Day	0
Rainfall	0
Sunshine	0
Humidity	0
Temperature	0

dtype: int64

Summary statistics:

	Year	Month	Day
Rainfall	Sunshine \		
count	76685.000000	76685.000000	76685.000000
	76685.000000		
mean	2020.499772	6.523962	15.726609
	5.914034		6.127835
std	1.707480	3.448442	8.798870
	3.269687		17.902041
min	2018.000000	1.000000	1.000000
	0.000000		0.000000
25%	2019.000000	4.000000	8.000000
	3.400000		0.000000
50%	2020.000000	7.000000	16.000000
	6.800000		0.000000
75%	2022.000000	10.000000	23.000000
	8.600000		2.000000
max	2023.000000	12.000000	31.000000
	28.500000		378.000000

	Humidity	Temperature
count	76685.000000	76685.000000
mean	79.458147	25.865309
std	8.355608	4.310779
min	32.000000	8.300000
25%	75.000000	22.800000
50%	80.000000	27.300000
75%	85.000000	29.200000
max	100.000000	34.300000

```
In [4]: # Feature engineering  
# Combine Year, Month, Day into datetime  
df['Date'] = pd.to_datetime(df[['Year','Month','Day']])  
df['DayOfYear'] = df['Date'].dt.dayofyear  
  
# Month as cyclical feature  
df['Month_sin'] = np.sin(2 * np.pi * df['Month']/12)  
df['Month_cos'] = np.cos(2 * np.pi * df['Month']/12)
```

```
In [5]: # Select features and target  
features = ['Humidity', 'Sunshine', 'Rainfall', 'DayOfYear',  
'Month_sin', 'Month_cos']  
X = df[features]  
y = df['Temperature']
```

```
In [6]: # Train-test split  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

```
In [7]: # Feature scaling for Linear Regression  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```


```
In [8]: # Linear Regression  
lr_model = LinearRegression()  
lr_model.fit(X_train_scaled, y_train)  
  
# Predictions  
y_train_pred_lr = lr_model.predict(X_train_scaled)  
y_test_pred_lr = lr_model.predict(X_test_scaled)
```

```
In [9]: # Metrics  
mae_train_lr = mean_absolute_error(y_train, y_train_pred_lr)  
mae_test_lr = mean_absolute_error(y_test, y_test_pred_lr)  
mse_train_lr = mean_squared_error(y_train, y_train_pred_lr)  
mse_test_lr = mean_squared_error(y_test, y_test_pred_lr)  
rmse_train_lr = np.sqrt(mse_train_lr)  
rmse_test_lr = np.sqrt(mse_test_lr)  
r2_train_lr = r2_score(y_train, y_train_pred_lr)  
r2_test_lr = r2_score(y_test, y_test_pred_lr)
```


```
In [10]: print("----- Linear Regression Metrics -----")
print(f"Train MAE: {mae_train_lr:.2f}, Test MAE:
{mae_test_lr:.2f}")
print(f"Train MSE: {mse_train_lr:.2f}, Test MSE:
{mse_test_lr:.2f}")
print(f"Train RMSE: {rmse_train_lr:.2f}, Test RMSE:
{rmse_test_lr:.2f}")
print(f"Train R2: {r2_train_lr:.2f}, Test R2: {r2_test_lr:.2f}")
```

```
----- Linear Regression Metrics -----
Train MAE: 1.75, Test MAE: 1.75
Train MSE: 4.82, Test MSE: 4.83
Train RMSE: 2.19, Test RMSE: 2.20
Train R2: 0.74, Test R2: 0.74
```

```
In [11]: # Residuals
residuals_lr = y_test - y_test_pred_lr
plt.figure(figsize=(8,5))
plt.scatter(y_test_pred_lr, residuals_lr, alpha=0.5, color='green')
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Predicted Temperature')
plt.ylabel('Residuals')
plt.title('Residual Plot (Linear Regression)')
plt.show()
```



```
In [12]: # Actual vs Predicted
plt.figure(figsize=(12,5))
plt.plot(y_test.values[:100], label='Actual', marker='o')
plt.plot(y_test_pred_lr[:100], label='Predicted', marker='x')
plt.title('Actual vs Predicted Temperature (Linear Regression, first 100 samples)')
plt.legend()
plt.show()
```



```
In [13]: # Ridge Regression
ridge_model = Ridge(alpha=1.0)
ridge_model.fit(X_train_scaled, y_train)


# Predictions
y_train_pred_ridge = ridge_model.predict(X_train_scaled)
y_test_pred_ridge = ridge_model.predict(X_test_scaled)
```

```
In [14]: # Metrics
mae_train_ridge = mean_absolute_error(y_train, y_train_pred_ridge)
mae_test_ridge = mean_absolute_error(y_test, y_test_pred_ridge)
mse_train_ridge = mean_squared_error(y_train, y_train_pred_ridge)
mse_test_ridge = mean_squared_error(y_test, y_test_pred_ridge)
rmse_train_ridge = np.sqrt(mse_train_ridge)
rmse_test_ridge = np.sqrt(mse_test_ridge)
r2_train_ridge = r2_score(y_train, y_train_pred_ridge)
r2_test_ridge = r2_score(y_test, y_test_pred_ridge)
```


```
In [15]: print("----- Ridge Regression Metrics -----")
print(f"Train MAE: {mae_train_ridge:.2f}, Test MAE:
{mae_test_ridge:.2f}")
print(f"Train MSE: {mse_train_ridge:.2f}, Test MSE:
{mse_test_ridge:.2f}")
print(f"Train RMSE: {rmse_train_ridge:.2f}, Test RMSE:
{rmse_test_ridge:.2f}")
print(f"Train R²: {r2_train_ridge:.2f}, Test R²:
{r2_test_ridge:.2f}")
```

```
----- Ridge Regression Metrics -----
Train MAE: 1.75, Test MAE: 1.75
Train MSE: 4.82, Test MSE: 4.83
Train RMSE: 2.19, Test RMSE: 2.20
Train R²: 0.74, Test R²: 0.74
```

```
In [16]: # Residuals
residuals_ridge = y_test - y_test_pred_ridge
plt.figure(figsize=(8,5))
plt.scatter(y_test_pred_ridge, residuals_ridge, alpha=0.5,
color='purple')
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Predicted Temperature')
plt.ylabel('Residuals')
plt.title('Residual Plot (Ridge Regression)')
plt.show()
```



```
In [17]: # Actual vs Predicted  
plt.figure(figsize=(12,5))  
plt.plot(y_test.values[:100], label='Actual', marker='o')  
plt.plot(y_test_pred_ridge[:100], label='Predicted', marker='x')  
plt.title('Actual vs Predicted Temperature (Ridge Regression, first 100 samples)')  
plt.legend()  
plt.show()
```




```
In [18]: # Lasso Regression  
lasso_model = Lasso(alpha=0.01)  
lasso_model.fit(X_train_scaled, y_train)  
  
# Predictions  
y_train_pred_lasso = lasso_model.predict(X_train_scaled)  
y_test_pred_lasso = lasso_model.predict(X_test_scaled)
```

```
In [19]: # Metrics  
mae_train_lasso = mean_absolute_error(y_train, y_train_pred_lasso)  
mae_test_lasso = mean_absolute_error(y_test, y_test_pred_lasso)  
mse_train_lasso = mean_squared_error(y_train, y_train_pred_lasso)  
mse_test_lasso = mean_squared_error(y_test, y_test_pred_lasso)  
rmse_train_lasso = np.sqrt(mse_train_lasso)  
rmse_test_lasso = np.sqrt(mse_test_lasso)  
r2_train_lasso = r2_score(y_train, y_train_pred_lasso)  
r2_test_lasso = r2_score(y_test, y_test_pred_lasso)
```


```
In [20]: print("----- Lasso Regression Metrics -----")
print(f"Train MAE: {mae_train_lasso:.2f}, Test MAE:
{mae_test_lasso:.2f}")
print(f"Train MSE: {mse_train_lasso:.2f}, Test MSE:
{mse_test_lasso:.2f}")
print(f"Train RMSE: {rmse_train_lasso:.2f}, Test RMSE:
{rmse_test_lasso:.2f}")
print(f"Train R^2: {r2_train_lasso:.2f}, Test R^2:
{r2_test_lasso:.2f}")
```

```
----- Lasso Regression Metrics -----
Train MAE: 1.75, Test MAE: 1.75
Train MSE: 4.82, Test MSE: 4.84
Train RMSE: 2.20, Test RMSE: 2.20
Train R^2: 0.74, Test R^2: 0.74
```

```
In [21]: # Residuals
residuals_lasso = y_test - y_test_pred_lasso
plt.figure(figsize=(8,5))
plt.scatter(y_test_pred_lasso, residuals_lasso, alpha=0.5,
color='orange')
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Predicted Temperature')
plt.ylabel('Residuals')
plt.title('Residual Plot (Lasso Regression)')
plt.show()
```



```
In [22]: # Actual vs Predicted
plt.figure(figsize=(12,5))
plt.plot(y_test.values[:100], label='Actual', marker='o')
plt.plot(y_test_pred_lasso[:100], label='Predicted', marker='x')
plt.title('Actual vs Predicted Temperature (Lasso Regression, first 100 samples)')
plt.legend()
plt.show()
```



```
In [23]: # Random Forest Regressor
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)


# Predictions
y_train_pred_rf = rf_model.predict(X_train)
y_test_pred_rf = rf_model.predict(X_test)
```

```
In [24]: # Metrics
mae_train_rf = mean_absolute_error(y_train, y_train_pred_rf)
mae_test_rf = mean_absolute_error(y_test, y_test_pred_rf)
mse_train_rf = mean_squared_error(y_train, y_train_pred_rf)
mse_test_rf = mean_squared_error(y_test, y_test_pred_rf)
rmse_train_rf = np.sqrt(mse_train_rf)
rmse_test_rf = np.sqrt(mse_test_rf)
r2_train_rf = r2_score(y_train, y_train_pred_rf)
r2_test_rf = r2_score(y_test, y_test_pred_rf)
```


```
In [25]: print("----- Random Forest Regressor Metrics -----")
print(f"Train MAE: {mae_train_rf:.2f}, Test MAE:
{mae_test_rf:.2f}")
print(f"Train MSE: {mse_train_rf:.2f}, Test MSE:
{mse_test_rf:.2f}")
print(f"Train RMSE: {rmse_train_rf:.2f}, Test RMSE:
{rmse_test_rf:.2f}")
print(f"Train R2: {r2_train_rf:.2f}, Test R2: {r2_test_rf:.2f}")
```

```
----- Random Forest Regressor Metrics -----
Train MAE: 0.40, Test MAE: 1.00
Train MSE: 0.32, Test MSE: 1.85
Train RMSE: 0.57, Test RMSE: 1.36
Train R2: 0.98, Test R2: 0.90
```

```
In [26]: # Residuals
residuals_rf = y_test - y_test_pred_rf
plt.figure(figsize=(8,5))
plt.scatter(y_test_pred_rf, residuals_rf, alpha=0.5, color='blue')
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Predicted Temperature')
plt.ylabel('Residuals')
plt.title('Residual Plot (Random Forest)')
plt.show()
```




```
In [27]: # Actual vs Predicted
plt.figure(figsize=(12,5))
plt.plot(y_test.values[:100], label='Actual', marker='o')
plt.plot(y_test_pred_rf[:100], label='Predicted', marker='x')
plt.title('Actual vs Predicted Temperature (Random Forest, first 100 samples)')
plt.legend()
plt.show()
```



```
In [28]: # Feature Importance
feature_importance = pd.DataFrame({
    'Feature': features,
    'Importance': rf_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

plt.figure(figsize=(8,5))
sns.barplot(x='Importance', y='Feature', data=feature_importance)
plt.title('Feature Importance (Random Forest)')
plt.show()
```




```
In [29]: print("\nRandom Forest Feature Importance:\n", feature_importance)
```


Random Forest Feature Importance:

	Feature	Importance
5	Month_cos	0.515236
4	Month_sin	0.200640
3	DayOfYear	0.165287
0	Humidity	0.058243
1	Sunshine	0.048348
2	Rainfall	0.012246


```
In [30]: # Correlation heatmap
plt.figure(figsize=(8,6))
sns.heatmap(df[['Temperature','Humidity','Sunshine','Rainfall']].corr(),
            annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Feature Correlation with Temperature')
plt.show()
```




```
In [31]: # Scatter plots for trend analysis  
plt.figure(figsize=(12,5))  
sns.scatterplot(x='Humidity', y='Temperature', data=df)  
plt.title('Temperature vs Humidity')  
plt.show()
```



```
In [32]: plt.figure(figsize=(12,5))  
sns.scatterplot(x='Sunshine', y='Temperature', data=df)  
plt.title('Temperature vs Sunshine')  
plt.show()
```



```
In [33]: plt.figure(figsize=(12,5))  
sns.scatterplot(x='Rainfall', y='Temperature', data=df)  
plt.title('Temperature vs Rainfall')  
plt.show()
```



Part–2: Classification → Rainfall Yes/No Prediction

```
In [34]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from imblearn.over_sampling import SMOTE
from sklearn.metrics import confusion_matrix, accuracy_score,
classification_report, roc_curve, auc
```

```
In [35]: # ----- CREATE BINARY TARGET -----
# 0 → No Rain, 1 → Rain
df['RainBinary'] = df['Rainfall'].apply(lambda x: 1 if x > 0 else
0)

print(df['RainBinary'].value_counts())
print(df['RainBinary'].unique())
```

```
RainBinary
0    52831
1    23854
Name: count, dtype: int64
[0 1]
```

```
In [36]: # ----- FEATURES -----
X = df[['Sunshine', 'Humidity', 'Temperature']]    # input
y = df['RainBinary']
```

```
In [37]: # ----- TRAIN TEST SPLIT -----
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
In [38]: # ----- SCALE (FOR LOGISTIC) -----
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [39]: # ----- SMOTE (FOR LR & RF) -----
sm = SMOTE(random_state=42)
```

```
In [40]: # LR SMOTE (scaled)
X_train_lr, y_train_lr = sm.fit_resample(X_train_scaled, y_train)

# RF SMOTE (non-scaled)
X_train_rf, y_train_rf = sm.fit_resample(X_train, y_train)

print("\nAfter SMOTE Rebalancing:", y_train_lr.value_counts())
```

After SMOTE Rebalancing: RainBinary
0 42265
1 42265
Name: count, dtype: int64

```
In [41]: # ----- LOGISTIC REGRESSION -----
log_model = LogisticRegression()
log_model.fit(X_train_lr, y_train_lr)

y_pred_log = log_model.predict(X_test_scaled)
y_prob_log = log_model.predict_proba(X_test_scaled)[:,1]
```

```
In [42]: print("\n===== Logistic Regression =====")
# Confusion Matrix
cm_log = confusion_matrix(y_test, y_pred_log)
print("Confusion Matrix:\n", cm_log)
# Accuracy
acc_log = accuracy_score(y_test, y_pred_log)
print(f"Accuracy: {acc_log:.2f}")
# Classification Report
print("\nClassification Report:\n", classification_report(y_test,
y_pred_log))
```


===== Logistic Regression =====
Confusion Matrix:
[[7627 2939]
 [912 3859]]
Accuracy: 0.75

	precision	recall	f1-score	support
0	0.89	0.72	0.80	10566
1	0.57	0.81	0.67	4771
accuracy			0.75	15337
macro avg	0.73	0.77	0.73	15337
weighted avg	0.79	0.75	0.76	15337

In [43]:

```
# ROC Curve
fpr_log, tpr_log, _ = roc_curve(y_test, y_prob_log)
roc_auc_log = auc(fpr_log, tpr_log)

plt.figure(figsize=(6,5))
plt.plot(fpr_log, tpr_log, label=f"AUC = {roc_auc_log:.2f}")
plt.plot([0,1],[0,1],'r--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve: Logistic Regression")
plt.legend()
plt.show()
```



In [44]:

```
# ----- RANDOM FOREST Classification -----
rf_model = RandomForestClassifier(n_estimators=200,
random_state=42)
rf_model.fit(X_train_rf, y_train_rf)

y_pred_rf = rf_model.predict(X_test)
y_prob_rf = rf_model.predict_proba(X_test)[:,1]
```

```
In [45]: print("\n===== Random Forest Classifier =====")  
cm_rf = confusion_matrix(y_test, y_pred_rf)  
print("Confusion Matrix:\n", cm_rf)  
acc_rf = accuracy_score(y_test, y_pred_rf)  
print(f"Accuracy: {acc_rf:.2f}")  
print("\nClassification Report:\n", classification_report(y_test,  
y_pred_rf))
```

===== Random Forest Classifier =====

Confusion Matrix:

```
[[8687 1879]  
[1753 3018]]
```


Accuracy: 0.76

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.82	0.83	10566
1	0.62	0.63	0.62	4771
accuracy			0.76	15337
macro avg	0.72	0.73	0.73	15337
weighted avg	0.76	0.76	0.76	15337

```
In [46]: fpr_rf, tpr_rf, _ = roc_curve(y_test, y_prob_rf)
roc_auc_rf = auc(fpr_rf, tpr_rf)

plt.figure(figsize=(6,5))
plt.plot(fpr_rf, tpr_rf, label=f"AUC = {roc_auc_rf:.2f}")
plt.plot([0,1],[0,1], 'r--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve: Random Forest")
plt.legend()
plt.show()
```



```
In [47]: !pip install xgboost
```

```
Requirement already satisfied: xgboost in
c:\users\hp\anaconda3\lib\site-packages (3.1.2)
Requirement already satisfied: numpy in
c:\users\hp\anaconda3\lib\site-packages (from xgboost)
(2.1.3)
Requirement already satisfied: scipy in
c:\users\hp\anaconda3\lib\site-packages (from xgboost)
(1.15.3)
```

```
In [48]: !pip install lightgbm
```

```
Requirement already satisfied: lightgbm in  
c:\users\hp\anaconda3\lib\site-packages (4.6.0)  
Requirement already satisfied: numpy>=1.17.0 in  
c:\users\hp\anaconda3\lib\site-packages (from lightgbm)  
(2.1.3)  
Requirement already satisfied: scipy in  
c:\users\hp\anaconda3\lib\site-packages (from lightgbm)  
(1.15.3)
```

```
In [49]: # Models  
from xgboost import XGBClassifier  
import xgboost as xgb  
import lightgbm as lgb
```

```
In [50]: # ----- XGBOOST (WITH CLASS WEIGHT) -----  
neg = (y_train == 0).sum()  
pos = (y_train == 1).sum()  
scale_pos_weight = neg / pos  
  
xgb_model = xgb.XGBClassifier(  
    eval_metric='logloss',  
    random_state=42,  
    scale_pos_weight=scale_pos_weight  
)  
xgb_model.fit(X_train, y_train)  
  
y_pred_xgb = xgb_model.predict(X_test)  
y_prob_xgb = xgb_model.predict_proba(X_test)[:,1]
```

```
In [51]: print("\n===== XGBoost Classifier =====")
cm_xgb = confusion_matrix(y_test, y_pred_xgb)
print("Confusion Matrix:\n", cm_xgb)
acc_xgb = accuracy_score(y_test, y_pred_xgb)
print(f"Accuracy: {acc_xgb:.2f}")
print("\nClassification Report:\n", classification_report(y_test,
y_pred_xgb))
```

===== XGBoost Classifier =====

Confusion Matrix:


```
[[7944 2622]
 [ 991 3780]]
```

Accuracy: 0.76

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.75	0.81	10566
1	0.59	0.79	0.68	4771
accuracy			0.76	15337
macro avg	0.74	0.77	0.75	15337
weighted avg	0.80	0.76	0.77	15337

```
In [52]: fpr_xgb, tpr_xgb, _ = roc_curve(y_test, y_prob_xgb)
roc_auc_xgb = auc(fpr_xgb, tpr_xgb)
plt.figure(figsize=(6,5))
plt.plot(fpr_xgb, tpr_xgb, label=f"AUC = {roc_auc_xgb:.2f}")
plt.plot([0,1],[0,1],'r--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve: XGBoost")
plt.legend()
plt.show()
```



```
In [53]: # ----- LIGHTGBM (WITH CLASS WEIGHT) -----
lgb_model = lgb.LGBMClassifier(
    random_state=42,
    is_unbalance=True
)
lgb_model.fit(X_train, y_train)

y_pred_lgb = lgb_model.predict(X_test)
y_prob_lgb = lgb_model.predict_proba(X_test)[:,1]
```


[LightGBM] [Info] Number of positive: 19083, number of negative: 42265
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.003581 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 422
[LightGBM] [Info] Number of data points in the train set: 61348, number of used features: 3
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.311061 -> initscore=-0.795161
[LightGBM] [Info] Start training from score -0.795161

```
In [54]: print("\n===== LightGBM Classifier =====")
cm_lgb = confusion_matrix(y_test, y_pred_lgb)
print("Confusion Matrix:\n", cm_lgb)
acc_lgb = accuracy_score(y_test, y_pred_lgb)
print(f"Accuracy: {acc_lgb:.2f}")
print("\nClassification Report:\n", classification_report(y_test, y_pred_lgb))
```

===== LightGBM Classifier =====
Confusion Matrix:
[[7933 2633]
 [959 3812]]
Accuracy: 0.77

	precision	recall	f1-score	support
0	0.89	0.75	0.82	10566
1	0.59	0.80	0.68	4771
accuracy			0.77	15337
macro avg	0.74	0.77	0.75	15337
weighted avg	0.80	0.77	0.77	15337

```
In [55]: fpr_lgb, tpr_lgb, _ = roc_curve(y_test, y_prob_lgb)
roc_auc_lgb = auc(fpr_lgb, tpr_lgb)
plt.figure(figsize=(6,5))
plt.plot(fpr_lgb, tpr_lgb, label=f"AUC = {roc_auc_lgb:.2f}")
plt.plot([0,1],[0,1],'r--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve: LightGBM")
plt.legend()
plt.show()
```



```
In [56]: # Accuracy comparison table
import pandas as pd


classifiers = ['Logistic Regression', 'Random Forest', 'XGBoost',
'LightGBM']
accuracies = [acc_log, acc_rf, acc_xgb, acc_lgb]
roc_aucs = [roc_auc_log, roc_auc_rf, roc_auc_xgb, roc_auc_lgb]

metrics_df = pd.DataFrame({'Classifier': classifiers, 'Accuracy': accuracies, 'ROC AUC': roc_aucs})
print("\nClassifier Performance:\n", metrics_df)

# Plot ROC Curves together
plt.figure(figsize=(8,6))
plt.plot(fpr_log, tpr_log, label=f'Logistic Regression (AUC={roc_auc_log:.2f})')
plt.plot(fpr_rf, tpr_rf, label=f'Random Forest (AUC={roc_auc_rf:.2f})')
plt.plot(fpr_xgb, tpr_xgb, label=f'XGBoost (AUC={roc_auc_xgb:.2f})')
plt.plot(fpr_lgb, tpr_lgb, label=f'LightGBM (AUC={roc_auc_lgb:.2f})')
plt.plot([0,1],[0,1], 'r--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curves Comparison")
plt.legend()
plt.show()
```

Classifier Performance:

	Classifier	Accuracy	ROC AUC
0	Logistic Regression	0.748908	0.847622
1	Random Forest	0.763187	0.825050
2	XGBoost	0.764426	0.861086
3	LightGBM	0.765795	0.863343



Part-3: Weather Station Similarity (Clustering)

In [57]: `from sklearn.cluster import KMeans`

```
In [58]: # Compute Average per Station
features = ['Temperature', 'Humidity', 'Rainfall', 'Sunshine']
df_avg = df.groupby('Station')[features].mean().reset_index()
print(df_avg.head(35))
```


	Station	Temperature	Humidity	Rainfall	Sunshine
0	Ambaganctg	26.302054	79.676403	7.695573	5.963350
1	Barisal	25.647056	82.578275	4.970790	6.175399
2	Bhola	25.919763	82.840256	5.370151	5.743131
3	Bogra	25.725057	77.068918	4.933364	5.386627
4	Chandpur	26.463943	79.556823	5.328161	5.725377
5	Chittagong	26.582793	77.563213	7.361935	6.279872
6	Chuadanga	25.158603	79.963943	3.783204	6.139890
7	Comilla	25.612095	79.719763	5.246006	6.209311
8	Coxsbazar	26.460292	78.158832	8.393884	6.959105
9	Dhaka	26.581105	72.482885	4.940666	5.155682
10	Dinajpur	25.049110	78.462346	5.265176	5.897718
11	Faridpur	25.989320	78.790050	5.183021	6.111456
12	Feni	25.510543	78.141944	7.042903	5.897627
13	Hatiya	25.901552	85.278412	8.093108	6.006664
14	Ishurdi	25.326654	78.826107	4.000456	6.215290
15	Jessore	26.019169	78.183934	4.081698	4.771383
16	Khepupara	26.111410	82.669101	8.006846	5.858330
17	Khulna	26.500365	78.657691	4.180739	6.308809
18	Kutubdia	26.418120	79.696942	7.605660	5.282930
19	Madaripur	25.825833	81.174806	5.056139	5.247604
20	Mcourt	26.349156	80.837974	6.668188	6.022866
21	Mongla	26.480100	80.981743	5.586947	5.839480
22	Mymensingh	25.209630	81.497033	5.731629	5.519306
23	Patuakhali	26.211319	84.189411	6.063441	6.294843
24	Rajshahi	25.487905	80.005933	3.357827	6.124236
25	Rangamati	25.627522	73.771337	6.582382	6.313692
26	Rangpur	25.177225	79.057052	5.780922	5.405340
27	Sandwip	26.131492	82.880420	8.542218	6.114925
28	Satkhira	26.374304	78.392971	4.092652	6.127111
29	Sitakunda	25.752351	75.913738	8.323140	6.117526
30	Srimangal	24.646417	81.575536	6.209037	6.095208
31	Sydpur	25.280420	77.772250	5.512551	5.633683
32	Sylhet	25.321543	78.248288	11.633044	5.456276
33	Tangail	25.632816	77.715199	5.087175	5.771474
34	Teknaf	26.498768	78.705614	8.763578	6.819671

```
In [59]: # Standardize Features
X = df_avg[features]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [60]: # not showing Windows+MKL warning
import warnings
warnings.filterwarnings("ignore", message=".*KMeans is known to
have a memory leak.*")
```


```
In [61]: # Elbow Method to find optimal k
inertia = []
K_range = range(1, 10)
for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_scaled)
    inertia.append(kmeans.inertia_)
```

```
In [62]: # Plot Elbow Curve
plt.figure(figsize=(8,5))
plt.plot(K_range, inertia, 'bo-', markersize=8)
plt.xlabel('Number of clusters (k)')
plt.ylabel('Inertia')
plt.title('Elbow Method for Optimal k')
plt.show()
```




```
In [63]: # Apply K-Means (Choose k=3 or 4)
k = 3 # or 4 based on elbow
kmeans = KMeans(n_clusters=k, random_state=42)
df_avg['Cluster'] = kmeans.fit_predict(X_scaled)
```

```
In [64]: # Cluster Plot (2D using Temperature & Humidity)
# indicate station Grouping
plt.figure(figsize=(8,6))
sns.scatterplot(x=df_avg['Temperature'], y=df_avg['Humidity'],
                 hue=df_avg['Cluster'], palette='Set2', s=100)
plt.title('Weather Stations Clustering')
plt.xlabel('Avg Temperature')
plt.ylabel('Avg Humidity')
plt.legend(title='Cluster')
plt.show()
```



```
In [65]: # Heatmap of clusters and features
cluster_means = df_avg.groupby('Cluster')[features].mean()
plt.figure(figsize=(8,6))
sns.heatmap(cluster_means, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Cluster Feature Averages')
plt.show()
```



```
In [66]: # Interpretation
for i in range(k):
    stations = df_avg[df_avg['Cluster'] == i]['Station'].tolist()
    print(f"\nCluster {i} Stations ({len(stations)}): {stations}")
```

Cluster 0 Stations (6): ['Bhola', 'Hatiya', 'Khepupara', 'Patuakhali', 'Sandwip', 'Sylhet']

Cluster 1 Stations (15): ['Barisal', 'Bogra', 'Chuadanga', 'Comilla', 'Dinajpur', 'Feni', 'Ishurdi', 'Jessore', 'Madaripur', 'Mymensingh', 'Rajshahi', 'Rangpur', 'Srimangal', 'Sydpur', 'Tangail']

Cluster 2 Stations (14): ['Ambaganctg', 'Chandpur', 'Chittagong', 'Coxsbazar', 'Dhaka', 'Faridpur', 'Khulna', 'Kutubdia', 'Mcourt', 'Mongla', 'Rangamati', 'Satkhira', 'Sitakunda', 'Teknaf']

```
In [67]: # Install plotly if not installed
# !pip install plotly

import plotly.express as px

# 3D interactive scatter plot
fig = px.scatter_3d(df_avg,
                     x='Temperature',
                     y='Humidity',
                     z='Rainfall',
                     color='Cluster',
                     hover_data=['Station'],
                     title='Interactive 3D Clustering of Weather
Stations',

color_continuous_scale=px.colors.qualitative.Set2,
                     height=700)

fig.update_traces(marker=dict(size=6))
fig.show()
```

```
In [68]: # Climate Zone
for c in range(k):
    print(f"\nCluster {c} Climate Pattern:")
    temp = cluster_means.loc[c, 'Temperature']
    hum = cluster_means.loc[c, 'Humidity']
    rain = cluster_means.loc[c, 'Rainfall']
    sun = cluster_means.loc[c, 'Sunshine']

    desc = ""

    if rain > cluster_means['Rainfall'].mean() and hum >
cluster_means['Humidity'].mean():
        desc = "Wet & Humid Zone"
    elif sun > cluster_means['Sunshine'].mean() and temp >
cluster_means['Temperature'].mean():
        desc = "Sunny & Hot Zone"
    elif temp < cluster_means['Temperature'].mean() and hum <
cluster_means['Humidity'].mean():
        desc = "Cool & Dry Zone"
    else:
        desc = "Mixed Climate Zone"

    print(desc)
```

Cluster 0 Climate Pattern:
Wet & Humid Zone

Cluster 1 Climate Pattern:
Cool & Dry Zone

Cluster 2 Climate Pattern:
Sunny & Hot Zone

```
In [ ]:
```